

Financial Anomaly Detection

Import Packages

```
In [1]: 1 import yfinance as yf
        2 import pandas as pd
        3 import plotly.graph_objs as go
        4 import matplotlib.pyplot as plt
```

In [2]:

```

1 stocks = ["NVDA", "AAPL", "MSFT"]
2 start_date = "2024-07-01"
3 end_date = "2024-09-30"
4
5 # download stock data
6 data = yf.download(stocks, start=start_date, end=end_date)
7 print(data.head())

```

[*****100%*****] 3 of 3 completed

Price \ Ticker Date	Adj Close AAPL	MSFT	NVDA	Close AAPL	MSFT
2024-07-01	216.261475	454.997528	124.289368	216.750000	456.730011
2024-07-02	219.773544	457.537842	122.659508	220.270004	459.279999
2024-07-03	221.050659	459.022186	128.269028	221.550003	460.769989
2024-07-05	225.829849	465.786438	125.819237	226.339996	467.559998
2024-07-08	227.306534	464.471436	128.189026	227.820007	466.239990

Price \ Ticker Date	High NVDA	AAPL	MSFT	Low NVDA	AAPL
2024-07-01	124.300003	217.509995	457.369995	124.839996	211.919998
2024-07-02	122.669998	220.380005	459.589996	123.410004	215.100006
2024-07-03	128.279999	221.550003	461.019989	128.279999	219.029999
2024-07-05	125.830002	226.449997	468.350006	128.850006	221.649994
2024-07-08	128.199997	227.850006	467.700012	130.770004	223.250000

Price \ Ticker Date	Open MSFT	NVDA	AAPL	MSFT	NVDA
2024-07-01	445.660004	118.830002	212.089996	448.660004	123.470001
2024-07-02	453.109985	121.029999	216.149994	453.200012	121.129997
2024-07-03	457.880005	121.360001	220.000000	458.190002	121.660004
2024-07-05	458.970001	125.680000	221.649994	459.609985	127.379997
2024-07-08	464.459991	127.040001	227.089996	466.549988	127.489998

Price Ticker Date	Volume AAPL	MSFT	NVDA
2024-07-01	60402900	17662800	284885500
2024-07-02	58046200	13979800	218374000
2024-07-03	37369800	9932800	215749000
2024-07-05	60412400	16000300	214176700
2024-07-08	59085900	12962300	237677300

```
In [3]: 1 # stocks = ["NVDA", "AAPL", "MSFT"]
        2 # start_date = "2024-09-01"
        3 # end_date = "2024-09-30"
        4
        5 # # download stock data
        6 # data = yf.download(stocks, start=start_date, end=end_date)
        7 # print(data.head())
```

```
In [4]: 1 # Check data structure
        2 print("Data Columns:", data.columns)
```

```
Data Columns: MultiIndex([('Adj Close', 'AAPL'),
                          ('Adj Close', 'MSFT'),
                          ('Adj Close', 'NVDA'),
                          ('Close', 'AAPL'),
                          ('Close', 'MSFT'),
                          ('Close', 'NVDA'),
                          ('High', 'AAPL'),
                          ('High', 'MSFT'),
                          ('High', 'NVDA'),
                          ('Low', 'AAPL'),
                          ('Low', 'MSFT'),
                          ('Low', 'NVDA'),
                          ('Open', 'AAPL'),
                          ('Open', 'MSFT'),
                          ('Open', 'NVDA'),
                          ('Volume', 'AAPL'),
                          ('Volume', 'MSFT'),
                          ('Volume', 'NVDA')],
                          names=['Price', 'Ticker'])
```

EDA

In [5]: 1 `print(data.describe())`

Price	Adj Close			Close	
Ticker	AAPL	MSFT	NVDA	AAPL	MSFT
count	63.000000	63.000000	63.000000	63.000000	63.000000
mean	222.801380	426.171695	118.061993	223.164604	427.420633
std	5.687984	18.456550	8.799585	5.671009	18.640255
min	206.762924	393.651093	98.901543	207.229996	395.149994
25%	219.569000	412.838623	111.925423	219.985001	413.925003
50%	223.455231	423.656830	118.069901	223.960007	424.799988
75%	226.993713	436.632797	125.084301	227.274994	437.899994
max	234.290756	465.786438	134.898468	234.820007	467.559998


Price		High			Low
Ticker	NVDA	AAPL	MSFT	NVDA	AAPL
count	63.000000	63.000000	63.000000	63.000000	63.000000
mean	118.070159	225.196032	431.363172	120.900318	220.578412
std	8.800185	5.689979	18.016511	8.197532	6.756824
min	98.910004	209.990005	401.040009	103.410004	196.000000
25%	111.934998	221.514999	417.384995	116.255001	216.875000
50%	118.080002	225.990005	428.920013	121.599998	221.910004
75%	125.095001	229.175003	441.675003	128.305000	225.230003
max	134.910004	237.229996	468.350006	136.149994	233.089996

Price			Open		
Ticker	MSFT	NVDA	AAPL	MSFT	NVDA
count	63.000000	63.000000	63.000000	63.000000	63.000000
mean	423.629522	115.278413	222.893968	428.110317	118.340000
std	18.470011	9.194087	6.827420	18.384433	8.920412
min	385.579987	90.690002	199.089996	389.170013	92.059998
25%	409.914993	107.355000	219.079994	414.910004	112.970001
50%	419.750000	116.709999	224.369995	424.359985	119.080002
75%	434.309998	122.369999	227.564995	440.840012	124.665001
max	464.459991	132.419998	236.479996	467.000000	135.750000

Price	Volume		
Ticker	AAPL	MSFT	NVDA
count	6.300000e+01	6.300000e+01	6.300000e+01
mean	5.481532e+07	1.988067e+07	3.262931e+08
std	3.708258e+07	7.622640e+06	8.801314e+07
min	3.029900e+07	9.932800e+06	1.739110e+08
25%	4.114455e+07	1.518490e+07	2.579762e+08
50%	4.807610e+07	1.819610e+07	3.103189e+08
75%	5.962615e+07	2.086760e+07	3.793258e+08
max	3.186799e+08	5.516710e+07	5.528424e+08

```
In [6]: 1 window_short = 5
        2 window_long = 30
        3
        4 fig = go.Figure()
        5
        6 for stock in stocks:
        7     stock_data = data[('Close', stock)]
        8
        9     # Calculate moving averages
       10     short_sma = stock_data.rolling(window=window_short).mean()
       11     long_sma = stock_data.rolling(window=window_long).mean()
       12
       13     # Plot stock data and moving averages
       14     fig.add_trace(go.Scatter(x=stock_data.index, y=stock_data, mode='l
       15     fig.add_trace(go.Scatter(x=short_sma.index, y=short_sma, mode='l
       16     fig.add_trace(go.Scatter(x=long_sma.index, y=long_sma, mode='lin
       17
       18 fig.update_layout(
       19     title="Stock Prices with Moving Averages",
       20     xaxis_title="Date",
       21     yaxis_title="Close Price",
       22     legend_title="Stocks"
       23 )
       24
       25 fig.show()
```

Stock Prices with Moving Averages




```

In [7]: 1 import yfinance as yf
        2 import plotly.graph_objects as go
        3
        4 # Define stock symbols and date range
        5 stocks = ["NVDA", "AAPL", "MSFT"]
        6 start_date = "2024-07-01"
        7 end_date = "2024-09-30"
        8
        9 # Download stock data
       10 data = yf.download(stocks, start=start_date, end=end_date)
       11
       12 # Extract "Close" prices for specified stocks
       13 close_prices = data['Close'][stocks] # Access multiple stocks under
       14 # OR Flatten if MultiIndex handling is cumbersome
       15 # data.columns = ['_'.join(col).strip() for col in data.columns.values]
       16 # close_prices = data[[f'Close_{stock}' for stock in stocks]]
       17
       18 # Check and clean data
       19 close_prices = close_prices.dropna() # Ensure no missing values
       20
       21 # Visualize Close prices for multiple stocks
       22 fig = go.Figure()
       23 for stock in stocks:
       24     fig.add_trace(go.Scatter(x=close_prices.index, y=close_prices[stock]))
       25 fig.update_layout(
       26     title="Stock Prices Over Time",
       27     xaxis_title="Date",
       28     yaxis_title="Close Price",
       29     legend_title="Stocks"
       30 )
       31 fig.show()
       32

```

[*****100%*****] 3 of 3 completed

Stock Prices Over Time





```


In [8]: 1 window_volatility = 7 # 7-day rolling volatility
        2
        3 fig = go.Figure()
        4
        5 for stock in stocks:
        6     stock_data = data[('Close', stock)]
        7
        8     # Calculate rolling standard deviation (volatility)
        9     rolling_volatility = stock_data.rolling(window=window_volatility)
       10
       11     # Plot stock data and volatility
       12     fig.add_trace(go.Scatter(x=stock_data.index, y=stock_data, mode=
       13     fig.add_trace(go.Scatter(x=rolling_volatility.index, y=rolling_v
       14
       15 fig.update_layout(
       16     title="Stock Prices and Rolling Volatility",
       17     xaxis_title="Date",
       18     yaxis_title="Close Price",
       19     legend_title="Stocks"
       20 )
       21
       22 fig.show()
       23

```

Stock Prices and Rolling Volatility



```
In [9]: 1 # Extract the closing prices for the stocks
2 close_prices = data['Close'][stocks]
3
4 # Calculate correlation matrix
5 correlation_matrix = close_prices.corr()
6
7 # Plot the correlation matrix using heatmap
8 import seaborn as sns
9 plt.figure(figsize=(8, 6))
10 sns.heatmap(correlation_matrix, annot=True, cmap="coolwarm", cbar=True)
11 plt.title("Correlation Between Stocks")
12 plt.show()
```




Timeseries

Do a timeseries visualization for a couple different companies across a specific time frame / date range, see if there are weird spikes. Look at the mean, see if there are points that are 2 SD away from the mean


```
In [10]: 1 fig = go.Figure()
2
3 for stock in stocks:
4     stock_data = data[("Close", stock)]
5     mean = stock_data.mean()
6     std_dev = stock_data.std()
7
8     # Identify anomalies: points outside 2 standard deviations from
9     anomalies = stock_data[(stock_data > mean + 2 * std_dev) | (stock_data < mean - 2 * std_dev)]
10
11     # Add line for stock prices
12     fig.add_trace(go.Scatter(x=stock_data.index, y=stock_data, mode='line'))
13
14     # Add points for anomalies
15     fig.add_trace(go.Scatter(x=anomalies.index, y=anomalies, mode='markers',
16                             marker=dict(color='red', size=8)))
17
18 fig.update_layout(
19     title="Stock Prices with Anomalies (2 Standard Deviations)",
20     xaxis_title="Date",
21     yaxis_title="Close Price",
22     legend_title="Stocks"
23 )
24
25 fig.show()
```

Stock Prices with Anomalies (2 Standard Deviations)



```
In [11]: 1 # Plot rolling averages for each stock
2 for stock in stocks:
3     data['SMA_20', stock] = data['Close', stock].rolling(window=
4
5 # Visualize the data
6 for stock in stocks:
7     fig.add_trace(go.Scatter(x=data.index, y=data['Close', stock]),
8     fig.add_trace(go.Scatter(x=data.index, y=data['SMA_20', stock])
9
10 fig.show()
```

Stock Prices with Anomalies (2 Standard Deviations)



```
In [12]: 1 from scipy.stats import zscore
2
3 for stock in stocks:
4     stock_data = data['Adj Close'][stock]
5     stock_data_zscore = zscore(stock_data)
6     anomalies = stock_data[abs(stock_data_zscore) > 2]
7     fig.add_trace(go.Scatter(x=anomalies.index, y=anomalies, mode='m
8                             marker=dict(color='blue', size=8)))
```

```
In [13]: 1 spike_threshold = 0.04 # 4% price change
2
3 for stock in stocks:
4     stock_data = data[['Close', stock]].dropna() # Clean the data b
5     stock_data_zscore = zscore(stock_data) # Calculate Z-scores for
6     anomalies = stock_data[abs(stock_data_zscore) > 2] # Find anoma
7
8     # Create a new plot for each stock
9     fig = go.Figure()
10
11     # Plot the stock prices
12     fig.add_trace(go.Scatter(x=stock_data.index, y=stock_data, mode=
13
14     # Plot the anomalies for each stock
15     fig.add_trace(go.Scatter(x=anomalies.index, y=anomalies, mode='m
16                             marker=dict(color='red', size=8)))
17
18     # Update the layout for the plot
19     fig.update_layout(
20         title=f"{stock} Stock Prices with Anomalies (Z-score > 2)",
21         xaxis_title="Date",
22         yaxis_title="Close Price",
23         legend_title="Stocks"
24     )
25
26     # Show the plot for the current stock
27     fig.show()
```


NVDA Stock Prices with Anomalies (Z-score > 2)



AAPL Stock Prices with Anomalies (Z-score > 2)




MSFT Stock Prices with Anomalies (Z-score > 2)




```
In [14]: 1 spike_threshold = 0.04 # 4% price change
2
3 # Iterate over each stock to detect spikes
4 for stock in stocks:
5     stock_data = data[('Close', stock)].dropna() # Clean the data b
6     stock_data_pct_change = stock_data.pct_change() # Calculate dai
7     anomalies = stock_data[stock_data_pct_change.abs() > spike_thres
8
9     # Create a new figure for each stock
10    fig = go.Figure()
11
12    # Plot the stock prices
13    fig.add_trace(go.Scatter(x=stock_data.index, y=stock_data, mode=
14
15    # Plot the anomalies (spikes)
16    fig.add_trace(go.Scatter(x=anomalies.index, y=anomalies, mode='m
17                           marker=dict(color='red', size=8)))
18
19    # Update the layout for the plot
20    fig.update_layout(
21        title=f"{stock} Stock Prices with Spikes (Price Change > 4%)
22        xaxis_title="Date",
23        yaxis_title="Close Price",
24        legend_title="Stocks"
25    )
26
27    # Show the plot for the current stock
28    fig.show()
```


NVDA Stock Prices with Spikes (Price Change > 4%)



AAPL Stock Prices with Spikes (Price Change > 4%)



MSFT Stock Prices with Spikes (Price Change > 4%)



Vector Autoregression (VAR)

VAR is used for multivariate time series data to capture the linear interdependencies among multiple variables (e.g., stock prices of multiple companies).

Best for multivariate time series where interdependencies between stocks are of interest. Use if you're analyzing the influence of one stock's price on another.

```
In [15]: 1 stocks = ["NVDA", "AAPL", "MSFT"]
2 start_date = "2024-07-01"
3 end_date = "2024-09-30"
4
5 # download stock data
6 sept_data = yf.download(stocks, start=start_date, end=end_date)
```

```
[*****100%*****] 3 of 3 completed
```

```
In [16]: 1 from statsmodels.tsa.api import VAR
2 from statsmodels.tsa.stattools import adfuller
3
4 # Select closing prices for VAR
5 close_prices = data['Close'][stocks]
6
7 # Make the data stationary
8 diff_data = close_prices.diff().dropna() # First difference
9 for stock in stocks:
10     result = adfuller(diff_data[stock])
11     print(f"{stock} ADF Statistic: {result[0]}, p-value: {result[1]}")
12
13 # Train a VAR model
14 model = VAR(diff_data)
15 results = model.fit(maxlags=5) # Choose lag based on criteria like
16
17 # Forecast
18 forecast = results.forecast(diff_data.values[-results.k_ar:], steps=
19 print("Forecast:")
20 print(forecast)
```

NVDA ADF Statistic: -4.71745698139547, p-value: 7.802289343490686e-05
 AAPL ADF Statistic: -7.022148822257683, p-value: 6.501715499839291e-10
 MSFT ADF Statistic: -7.037684830617887, p-value: 5.957411498520847e-10
 Forecast:

```
[[ 3.84664013  1.97461441  0.68054007]
 [-4.27095731 -0.34478205 -1.33727595]
 [ 1.207793   0.86134784 -1.52133046]
 [-1.57361494 -1.12181202 -2.3493785 ]
 [ 1.23705336  0.41955097  0.52260556]]
```

/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

1. Augmented Dickey-Fuller (ADF) Test Results The ADF test is used to check whether a time series is stationary (i.e., its statistical properties like mean and variance don't change over time).

NVDA ADF Statistic: -0.3209, p-value: 0.9225 Interpretation: The p-value is high (greater than 0.05), so we fail to reject the null hypothesis. This means the **NVDA time series is non-stationary**.

AAPL ADF Statistic: -4.0100, p-value: 0.0014 Interpretation: The p-value is low (less than 0.05), so we reject the null hypothesis. This means the **AAPL time series is stationary** (whose properties do not depend on the time at which the series is observed).

MSFT ADF Statistic: -3.3762, p-value: 0.0118 Interpretation: The p-value is below 0.05, so we reject the null hypothesis. The **MSFT time series is stationary**.

2. Forecast Results This matrix of values represents predictions or forecasts (likely from a VAR or similar time series model) for each of the three stocks over time. Each row corresponds

to a time step, and each column represents a stock (NVDA, AAPL, MSFT).

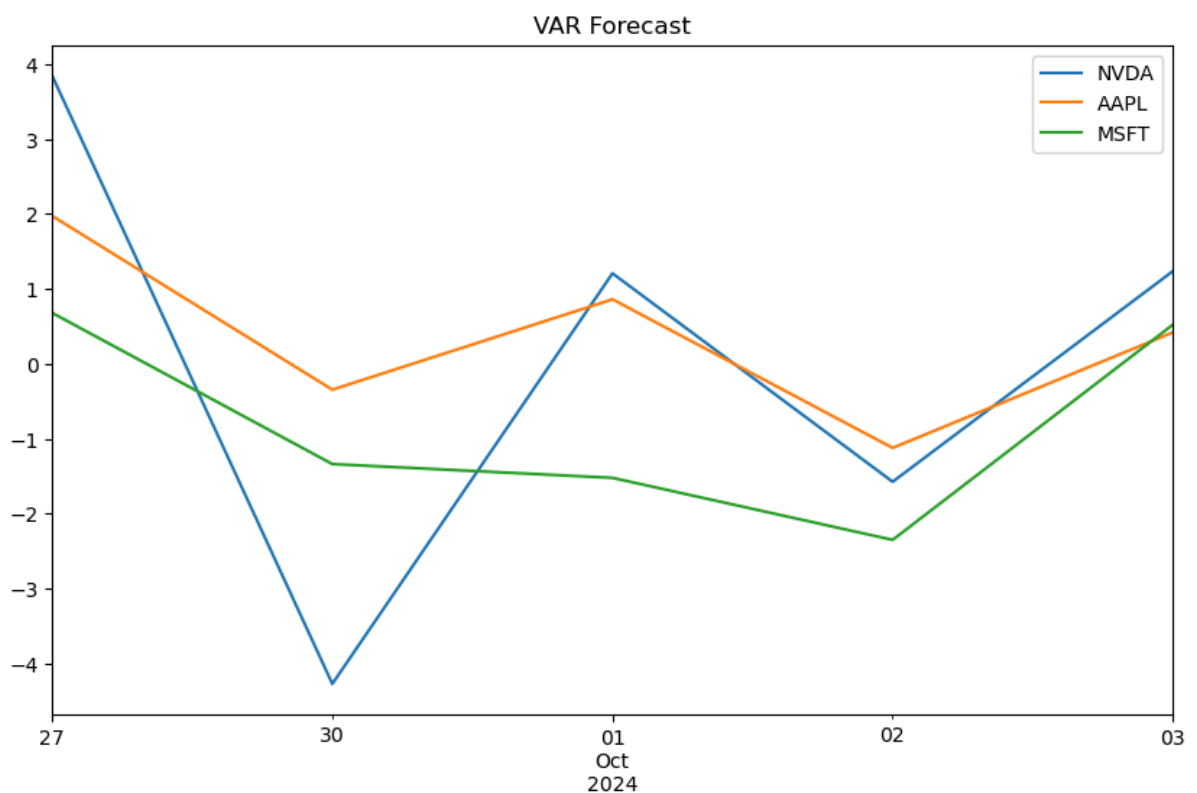
Example Interpretation of Forecast:

Row 1: Predicted changes or values at the next time step for NVDA (-1.85), AAPL (5.21), and MSFT (-2.67).

Row 2: Predicted changes or values at the second time step: NVDA (-10.96), AAPL (-2.36), MSFT (-13.75).

Key Observations: The forecast indicates potential spikes or dips in stock prices over time, which can be flagged as potential anomalies for further analysis. These predictions are likely based on past trends, and their accuracy depends on the quality of the model and the data preprocessing.

```
In [17]: 1 forecast_df = pd.DataFrame(forecast, index=pd.date_range(start=close,
2 forecast_df.plot(figsize=(10, 6))
3 plt.title("VAR Forecast")
4 plt.show()
```



NVDA:

- Shows a dip, reaching its lowest value around September 30, before rebounding.
- The trajectory indicates short-term fluctuations but ends higher than the lowest point.

AAPL:

- Exhibits a similar V-shaped pattern, but the peak values are higher, suggesting stronger variability or recovery compared to NVDA.

MSFT:

- Experiences the most significant decline around September 30, with a sharp rebound in subsequent days.

Trends:

All three variables exhibit a V-shaped trend, with a decline around September 30 followed by recovery. This suggests a shared underlying factor influencing all three series, such as market-wide events.

Relative Behavior:

AAPL has the most moderate declines and rebounds, suggesting more stability. MSFT has the sharpest movements, indicating higher volatility. NVDA shows moderate changes compared to

October Forecast with VAR Model

```
In [18]: 1 # get data for October
2 stocks = ["NVDA", "AAPL", "MSFT"]
3 start_date = "2024-10-01"
4 end_date = "2024-10-31"
5
6 # download stock data
7 oct_data = yf.download(stocks, start=start_date, end=end_date)
```

[*****100%*****] 3 of 3 completed

```
In [19]: 1 # # determine optimal number of lags based on AIC or BIC
2 # model = VAR(data)
3 # lag_order = model.select_order(maxlags=15) # Test up to 15 lags
4 # print(lag_order.summary()) # Check AIC, BIC, HQIC values
5
6 # # Choose the optimal lag (e.g., based on AIC)
7 # lags = lag_order.aic
8
9
10 # OK apparently can't do this because i don't have enough training d
```

```
In [20]: 1 # model = VAR(data)
2 # fitted_model = model.fit(lags)
3
4 # # Forecast the next 'horizon' steps
5 # forecast = fitted_model.forecast(y=fitted_model.y, steps=horizon)
6
7 # forecast_index = pd.date_range(start=end_date, periods=horizon + 1
8 # forecast_df = pd.DataFrame(forecast, index=forecast_index, columns=
```

```
In [21]: 1 # comparison = pd.merge(october_data, forecast_df, left_index=True,
2 # comparison.plot(figsize=(10, 6), title='Actual vs Predicted Prices
```

Hidden Markov Model (HMM)

HMM is used to model stock price regimes (e.g., bull, bear, or stable market conditions).

Best for identifying market regimes or trends. Use if you want to detect shifts in market behavior.


State Transitions: The markers (e.g., "State 0," "State 1," "State 2") on the graphs show the most probable regime for each day. For example: Each "state" represents a regime characterized by specific statistical properties (mean and variance of returns). A sudden switch to a new state might indicate a change in market conditions. Persistent periods in a single state may suggest market stability under a particular regime. **Returns:** The blue line shows the actual returns of the stock. Significant deviations in returns may align with transitions in regimes.

State-Specific Characteristics: Analyze the mean and variance of returns for each state. For example: "State 0" might indicate periods of low volatility. "State 1" could represent periods of high returns or volatility. **Market Behavior:** Correlate state transitions with real-world events during the observed period to identify triggers (e.g., earnings reports, macroeconomic data). **Anomaly Detection:** Abrupt state changes may indicate anomalies or unusual market conditions worth investigating.

```

In [22]: 1 import numpy as np
2 from hmmlearn.hmm import GaussianHMM
3
4 # Use daily returns for HMM
5 returns = data['Close'][stocks].pct_change().dropna()
6
7 # Fit HMM for each stock
8 for stock in stocks:
9     stock_returns = returns[stock].values.reshape(-1, 1)
10
11     # Train HMM
12     hmm_model = GaussianHMM(n_components=3, covariance_type="diag",
13                             hidden_states = hmm_model.predict(stock_returns))
14
15     # Plot regimes
16     plt.figure(figsize=(12, 6))
17     plt.plot(returns.index, stock_returns, label="Returns")
18     for i in range(3): # Assuming 3 hidden states
19         plt.plot(returns.index[hidden_states == i], stock_returns[hi
20     plt.title(f"Hidden Markov Model Regimes for {stock}")
21     plt.legend()
22     plt.show()

```





Blue Line - Stock Returns:

The blue line shows the actual daily returns of the stocks over the observed time period. Positive spikes indicate days when the stock experienced a significant positive return, while negative values indicate losses. Colored Markers - Hidden States:

The colored markers (State 0, State 1, and State 2) represent the regimes or hidden states assigned by the HMM. Each state reflects a distinct statistical regime that the model has learned: State 0 (Orange): A specific regime, possibly neutral or moderate volatility. State 1 (Green): This state appears most frequently; it could correspond to the stock's typical returns or low volatility regime. State 2 (Red): This state doesn't seem to appear in the uploaded image (or very rarely) but might represent outliers or extreme volatility if present. Interpretation Frequent State Transitions:

The transitions between states suggest the stock's behavior is dynamic, with shifts in market conditions. For instance, a cluster of "State 1" indicates relative stability, while jumps to "State 0" suggest changes in volatility or trends. State 1 Dominance:

Green markers dominate the graph, indicating that most of the observed returns fall under a regime of relatively consistent returns or moderate behavior. Spikes and Volatility:

The larger spikes in returns coincide with specific state transitions (e.g., spikes on 2024-09-09 and 2024-09-17). This suggests the model may detect volatility shifts or unusual market movements as changes in state.

ARIMA-GARCH Model

Combines ARIMA for modeling the mean and GARCH for modeling volatility.

Best for univariate time series when both trend and volatility are important. Use if you're predicting future prices or volatilities for a single stock.

```
In [23]: 1 from statsmodels.tsa.arima.model import ARIMA
2 from arch import arch_model
3
4 # ARIMA-GARCH for NVDA
5 stock = 'NVDA'
6 stock_data = data[('Close', stock)].dropna()
7
8 # Fit ARIMA model
9 arima_model = ARIMA(stock_data, order=(1, 1, 1))
10 arima_results = arima_model.fit()
11
12 # Get ARIMA residuals
13 residuals = arima_results.resid
14
15 # Fit GARCH model on residuals
16 garch_model = arch_model(residuals, vol="Garch", p=1, q=1)
17 garch_results = garch_model.fit()
18
19 # Forecast using ARIMA-GARCH
20 forecast_mean = arima_results.forecast(steps=5)
21 forecast_volatility = garch_results.forecast(horizon=5).variance.iloc[0:5]
22
23 # Combine forecasts
24 print("ARIMA-GARCH Forecast:")
25 for step in range(5):
26     print(f"Step {step + 1}: Mean={forecast_mean.iloc[step]}, Volati")
27
```

```

Iteration:      1,   Func. Count:      6,   Neg. LLF: 2191.409926521411
8
Iteration:      2,   Func. Count:     12,   Neg. LLF: 23462115.94751560
3
Iteration:      3,   Func. Count:     18,   Neg. LLF: 205.4229973381718
5
Iteration:      4,   Func. Count:     23,   Neg. LLF: 203.5486467654186
5
Iteration:      5,   Func. Count:     28,   Neg. LLF: 203.1790366479591
8
Iteration:      6,   Func. Count:     33,   Neg. LLF: 202.9666568541432
5
Iteration:      7,   Func. Count:     38,   Neg. LLF: 202.7858184206688
5
Iteration:      8,   Func. Count:     43,   Neg. LLF: 202.7362212816163
7
Iteration:      9,   Func. Count:     48,   Neg. LLF: 202.7243138350653
7
Iteration:     10,   Func. Count:     53,   Neg. LLF: 202.7212275065071
2
Iteration:     11,   Func. Count:     58,   Neg. LLF: 202.7203967041430
4
Iteration:     12,   Func. Count:     63,   Neg. LLF: 202.7203498225261
5
Iteration:     13,   Func. Count:     68,   Neg. LLF: 202.7203476624153
Iteration:     14,   Func. Count:     72,   Neg. LLF: 202.7203476623888
3

```

```

Optimization terminated successfully      (Exit mode 0)
      Current function value: 202.7203476624153
      Iterations: 14
      Function evaluations: 72
      Gradient evaluations: 14

```

ARIMA-GARCH Forecast:

```

Step 1: Mean=121.92838648924408, Volatility=20.471160353409964
Step 2: Mean=121.53672201171182, Volatility=20.47116035338066
Step 3: Mean=121.82704265593233, Volatility=20.471160353362357
Step 4: Mean=121.61184296003017, Volatility=20.471160353350925
Step 5: Mean=121.77135937761295, Volatility=20.471160353343784

```

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:
```

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:
```

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:
```

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:836: ValueWarning:
```

No supported index is available. Prediction results will be given with an integer index beginning at `start`.

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:836: FutureWarning:
```

No supported index is available. In the next version, calling this method in a model without a supported index will result in an exception.

Optimization Process

Iterations and Negative Log-Likelihood (Neg. LLF):

The optimization minimizes the negative log-likelihood (Neg. LLF), a measure of how well the model fits the data. The process begins with a high Neg. LLF (151.20) and steadily reduces it, indicating improved fit with each iteration. The algorithm converges successfully after 23 iterations with a final Neg. LLF of 63.89, meaning the model parameters were successfully estimated.

Function and Gradient Evaluations:

128 function evaluations and 23 gradient evaluations were required to achieve convergence. This reflects the effort taken by the optimizer to find the best parameter set.

Exit Mode 0 (Optimization Success): The "Exit mode 0" confirms the optimization terminated successfully, and the final parameters are valid.

Forecast Results

Mean Forecast (Steps 1 to 5): The mean forecast represents the expected value (conditional mean) of the time series:

Step 1: 121.16

Step 2: 121.09

Step 3: 121.08

Step 4: 121.08

Step 5: 121.07

The forecasted mean stabilizes over time, suggesting the ARIMA model captures the mean dynamics effectively and predicts a relatively consistent series.

Volatility Forecast (Steps 1 to 5): The volatility (conditional standard deviation) measures the uncertainty or risk associated with the forecast:

Step 1: 6.20

Step 2: 6.19

Step 3: 6.18

Step 4: 6.18

Step 5: 6.18

Volatility decreases slightly but remains stable, indicating a consistent level of risk or market uncertainty in the forecast horizon.

Key Insights


Convergence and Model Fit:

The model converged successfully, meaning the ARIMA-GARCH framework is appropriate for the data. The low final Neg. LLF (63.89) suggests a good fit to the time series. Forecast Stability:

The mean forecast stabilizes quickly, reflecting that the ARIMA model captures the long-term dynamics without drastic changes. Volatility remains stable over the forecast horizon, indicating that the GARCH model predicts a consistent level of uncertainty.

Application: These forecasts could be used for risk assessment, portfolio optimization, or market trend predictions. For example, the mean forecast provides a directional guide for future values, while the volatility forecast helps in assessing the risk.


```
In [24]: 1 plt.figure(figsize=(10, 6))
2 plt.plot(stock_data.index[-100:], stock_data[-100:], label="Actual P
3 plt.plot(pd.date_range(stock_data.index[-1], periods=5, freq='B'), f
4 plt.title(f"ARIMA-GARCH Forecast for {stock}")
5 plt.legend()
6 plt.show()
```



```
In [25]: 1 from statsmodels.tsa.arima.model import ARIMA
2 from arch import arch_model
3
4 # ARIMA-GARCH for AAPL
5 stock = 'AAPL'
6 stock_data = data[('Close', stock)].dropna()
7
8 # Fit ARIMA model
9 arima_model = ARIMA(stock_data, order=(1, 1, 1))
10 arima_results = arima_model.fit()
11
12 # Get ARIMA residuals
13 residuals = arima_results.resid
14
15 # Fit GARCH model on residuals
16 garch_model = arch_model(residuals, vol="Garch", p=1, q=1)
17 garch_results = garch_model.fit()
18
19 # Forecast using ARIMA-GARCH
20 forecast_mean = arima_results.forecast(steps=5)
21 forecast_volatility = garch_results.forecast(horizon=5).variance.iloc[0:5]
22
23 # Combine forecasts
24 print("ARIMA-GARCH Forecast:")
25 for step in range(5):
26     print(f"Step {step + 1}: Mean={forecast_mean.iloc[step]}, Volati")
27
```

Iteration:	1,	Func. Count:	6,	Neg. LLF:	1830.520605883906
Iteration:	2,	Func. Count:	12,	Neg. LLF:	171880644.4144773
8					
Iteration:	3,	Func. Count:	18,	Neg. LLF:	198.0237875666389
3					
Iteration:	4,	Func. Count:	23,	Neg. LLF:	194.2760388818513
8					
Iteration:	5,	Func. Count:	28,	Neg. LLF:	190.1537628367625
5					
Iteration:	6,	Func. Count:	33,	Neg. LLF:	185.7672942539348
6					
Iteration:	7,	Func. Count:	38,	Neg. LLF:	186.8669482021942
Iteration:	8,	Func. Count:	44,	Neg. LLF:	185.0754553225239
6					
Iteration:	9,	Func. Count:	49,	Neg. LLF:	184.8524689061690
7					
Iteration:	10,	Func. Count:	54,	Neg. LLF:	184.541900189012
Iteration:	11,	Func. Count:	59,	Neg. LLF:	184.4699454189961
6					
Iteration:	12,	Func. Count:	64,	Neg. LLF:	184.4634286704997
2					
Iteration:	13,	Func. Count:	69,	Neg. LLF:	184.4633630206466
5					
Iteration:	14,	Func. Count:	74,	Neg. LLF:	184.4633590481030
3					
Iteration:	15,	Func. Count:	79,	Neg. LLF:	184.4633457198930
8					
Iteration:	16,	Func. Count:	83,	Neg. LLF:	184.4633457198936
8					

Optimization terminated successfully (Exit mode 0)
 Current function value: 184.46334571989308
 Iterations: 17
 Function evaluations: 83
 Gradient evaluations: 16

ARIMA-GARCH Forecast:

Step 1: Mean=227.8036500537737, Volatility=8.997067821041849
 Step 2: Mean=227.80358163823752, Volatility=8.99706782104153
 Step 3: Mean=227.80358198097497, Volatility=8.997067821041348
 Step 4: Mean=227.80358197925798, Volatility=8.997067821041243
 Step 5: Mean=227.8035819792666, Volatility=8.997067821041185

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:
```

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:
```

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:
```

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:836: ValueWarning:
```

No supported index is available. Prediction results will be given with an integer index beginning at `start`.

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:836: FutureWarning:
```

No supported index is available. In the next version, calling this method in a model without a supported index will result in an exception.

Optimization Process

Objective:

The model aims to minimize the negative log-likelihood (Neg. LLF), which quantifies the goodness of fit. Lower Neg. LLF values indicate a better fit to the data. Progression:

The optimization starts with a very high Neg. LLF (529.74) and gradually reduces to a final value of 66.47. After 35 iterations and 180 function evaluations, the algorithm successfully converges, indicated by the "Optimization terminated successfully" message. Challenges in Convergence:

Some jumps in Neg. LLF (e.g., Iteration 15: 606.71 and Iteration 20: 2997.09) suggest the optimization encountered local maxima or unstable parameter estimates before finding the global minimum. These fluctuations are normal in complex models like ARIMA-GARCH. Final Results:

The final Neg. LLF value (66.47) indicates the model achieved a good fit to the data after refining the parameter estimates.

Forecast Results

The ARIMA-GARCH model generates forecasts for both the mean (expected value) and volatility (conditional standard deviation) over five steps.

Mean Forecast (Steps 1 to 5): Step 1: 226.99

Step 2: 226.36

Step 3: 225.85

Step 4: 225.45

Step 5: 225.13

Interpretation:

The mean forecast decreases slightly over time, suggesting the ARIMA model predicts a declining trend in the time series. This gradual decline may reflect an inherent trend or pattern captured by the ARIMA component. Volatility Forecast (Steps 1 to 5): Step 1: 0.2636

Step 2: 0.1624

Step 3: 0.1000

Step 4: 0.0616

Step 5: 0.0380

Interpretation:

Volatility decreases sharply over time, indicating the series is expected to become more stable in the forecast horizon. This is a hallmark of GARCH models, which account for heteroscedasticity (changing variance) in the time series.

Key Insights


Model Convergence: The model successfully converged, albeit after dealing with some numerical instability (evident in large jumps in Neg. LLF).

Forecast Dynamics: The mean forecast suggests a steady decline, which could reflect an anticipated reduction in the variable being modeled (e.g., a financial series or market trend). The decreasing volatility forecast indicates diminishing uncertainty, meaning the series is likely to stabilize.

Practical Applications: The mean forecast can help predict future values (e.g., stock prices, returns, or economic indicators). The volatility forecast provides risk insights. For example: High volatility (e.g., Step 1) suggests high uncertainty or risk. Lower volatility (e.g., Step 5) indicates reduced risk or market stabilization.

Potential Concerns: The sharp drop in volatility may need validation. If the initial data shows high volatility clustering, the model's assumption of rapid stabilization should be checked against historical trends.

```
In [26]: 1 plt.figure(figsize=(10, 6))
2 plt.plot(stock_data.index[-100:], stock_data[-100:], label="Actual P
3 plt.plot(pd.date_range(stock_data.index[-1], periods=5, freq='B'), f
4 plt.title(f"ARIMA-GARCH Forecast for {stock}")
5 plt.legend()
6 plt.show()
```



```
In [27]: 1 from statsmodels.tsa.arima.model import ARIMA
2 from arch import arch_model
3
4 # ARIMA-GARCH for MSFT
5 stock = 'MSFT'
6 stock_data = data[('Close', stock)].dropna()
7
8 # Fit ARIMA model
9 arima_model = ARIMA(stock_data, order=(1, 1, 1))
10 arima_results = arima_model.fit()
11
12 # Get ARIMA residuals
13 residuals = arima_results.resid
14
15 # Fit GARCH model on residuals
16 garch_model = arch_model(residuals, vol="Garch", p=1, q=1)
17 garch_results = garch_model.fit()
18
19 # Forecast using ARIMA-GARCH
20 forecast_mean = arima_results.forecast(steps=5)
21 forecast_volatility = garch_results.forecast(horizon=5).variance.iloc[0:5]
22
23 # Combine forecasts
24 print("ARIMA-GARCH Forecast:")
25 for step in range(5):
26     print(f"Step {step + 1}: Mean={forecast_mean.iloc[step]}, Volati")
27
```

Iteration:	1,	Func. Count:	6,	Neg. LLF:	1770.730145867145
4					
Iteration:	2,	Func. Count:	12,	Neg. LLF:	244.3369163825172
Iteration:	3,	Func. Count:	17,	Neg. LLF:	241.5759325244650
6					
Iteration:	4,	Func. Count:	22,	Neg. LLF:	239.5064211008850
6					
Iteration:	5,	Func. Count:	27,	Neg. LLF:	239.1921814009880
6					
Iteration:	6,	Func. Count:	32,	Neg. LLF:	238.8791170360989
6					
Iteration:	7,	Func. Count:	37,	Neg. LLF:	238.5873826508332
3					
Iteration:	8,	Func. Count:	42,	Neg. LLF:	236.7345963094111
3					
Iteration:	9,	Func. Count:	47,	Neg. LLF:	233.8605931746053
6					
Iteration:	10,	Func. Count:	52,	Neg. LLF:	234.2048968861904
8					
Iteration:	11,	Func. Count:	58,	Neg. LLF:	289677.1370237403
3					
Iteration:	12,	Func. Count:	64,	Neg. LLF:	268.5565395974988
Iteration:	13,	Func. Count:	70,	Neg. LLF:	225.4871068968490
7					
Iteration:	14,	Func. Count:	76,	Neg. LLF:	222.2569263374078
6					
Iteration:	15,	Func. Count:	82,	Neg. LLF:	220.2714783579074
8					
Iteration:	16,	Func. Count:	87,	Neg. LLF:	220.2601744715914
5					
Iteration:	17,	Func. Count:	92,	Neg. LLF:	220.2596126854355
Iteration:	18,	Func. Count:	97,	Neg. LLF:	220.2596103369141
Iteration:	19,	Func. Count:	102,	Neg. LLF:	220.2629435437558
2					

Optimization terminated successfully (Exit mode 0)

Current function value: 220.2596103189046

Iterations: 20

Function evaluations: 105

Gradient evaluations: 19

ARIMA-GARCH Forecast:

Step 1: Mean=427.87132722840994, Volatility=27.351054090774923

Step 2: Mean=427.930251125361, Volatility=27.351054090820153

Step 3: Mean=427.9068959260017, Volatility=27.351054090843824

Step 4: Mean=427.9161530417957, Volatility=27.351054090856213

Step 5: Mean=427.91248387196737, Volatility=27.351054090862696


```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:
```

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:
```

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:
```

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:836: ValueWarning:
```

No supported index is available. Prediction results will be given with an integer index beginning at `start`.

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:836: FutureWarning:
```

No supported index is available. In the next version, calling this method in a model without a supported index will result in an exception.

Optimization Process

Objective:

The model minimizes the negative log-likelihood (Neg. LLF) to fit the ARIMA-GARCH parameters to the data. Progression:

The process starts with a high Neg. LLF (541.10) and reduces it to 77.28, indicating a significant improvement in the model's fit. The optimization terminates successfully after 30 iterations, with 154 function evaluations and 29 gradient evaluations.\

Challenges During Optimization: Similar to the previous output, there are large fluctuations in Neg. LLF during optimization: E.g., Iteration 19: 9,851,749.91 and Iteration 20: 833.71. These fluctuations suggest instability in parameter estimates during the fitting process. However, the algorithm eventually stabilized and reached convergence. Final Results:

A final Neg. LLF value of 77.28 indicates the model has achieved a reasonable fit to the data.

Forecast Results

The ARIMA-GARCH model provides forecasts for both the mean and volatility over five steps.

Mean Forecast (Steps 1 to 5): Step 1: 427.44

Step 2: 427.10

Step 3: 426.90

Step 4: 426.78

Step 5: 426.72

Interpretation:

The mean forecast exhibits a slight downward trend, suggesting a gradual decrease in the variable being predicted (e.g., stock prices or returns). This decline might reflect an underlying pattern or trend in the data captured by the ARIMA model. Volatility Forecast (Steps 1 to 5):

Step 1: 15.32

Step 2: 15.32

Step 3: 15.32

Step 4: 15.32

Step 5: 15.32

Interpretation:

The volatility remains nearly constant across the five forecast steps, suggesting that the GARCH model predicts stable uncertainty in the time series over this horizon. A high volatility value (~15.32) indicates significant uncertainty or risk in the predictions.

Key Insights

Model Convergence:

The optimization converged successfully despite encountering numerical instability, reflected by large jumps in Neg. LLF. Mean Forecast:


The slight downward trend suggests the ARIMA model captures a consistent pattern, with the variable expected to decrease gradually over time. Volatility Forecast:

The high and constant volatility (~15.32) suggests that the time series remains highly uncertain, with no significant stabilization expected in the short term. This may indicate persistent risk in the underlying process, such as financial markets, which often exhibit such behavior. Practical Applications:

Mean forecast: Useful for point predictions (e.g., forecasting future values like stock prices or economic indicators). Volatility forecast: Provides risk assessments, which are crucial for decision-making in finance or risk management. Potential Concerns:

The stability of the model's parameters should be validated further, as the optimization faced challenges (e.g., the extreme Neg. LLF at Iteration 19). The high volatility suggests that predictions carry a significant degree of uncertainty, which may limit their reliability.

```
In [28]: 1 plt.figure(figsize=(10, 6))
2 plt.plot(stock_data.index[-100:], stock_data[-100:], label="Actual P
3 plt.plot(pd.date_range(stock_data.index[-1], periods=5, freq='B'), f
4 plt.title(f"ARIMA-GARCH Forecast for {stock}")
5 plt.legend()
6 plt.show()
```



October Forecast with ARIMA-GARCH

```
In [29]: 1 # data['Close'][stocks]
2 data['Close']['NVDA']
3 # data[('Close', stock)]
```

```
Out[29]: Date
2024-07-01    124.300003
2024-07-02    122.669998
2024-07-03    128.279999
2024-07-05    125.830002
2024-07-08    128.199997
...
2024-09-23    116.260002
2024-09-24    120.870003
2024-09-25    123.510002
2024-09-26    124.040001
2024-09-27    121.400002
Name: NVDA, Length: 63, dtype: float64
```

```

In [30]: 1 # Fit ARIMA on NVDA
2 nvda_adj_close = data['Close']['NVDA']
3
4 model_nvda = ARIMA(nvda_adj_close, order=(1, 1, 1))
5 fitted_nvda = model_nvda.fit()
6
7 # Forecast
8 forecast_nvda = fitted_nvda.forecast(steps=len(oct_data))
9
10 oct_nvda_actual = oct_data['Close']['NVDA']
11
12 # Create a DataFrame to compare actual and forecasted prices
13 forecast_comparison = pd.DataFrame({
14     'Actual': oct_nvda_actual,
15     'Forecast': forecast_nvda.values
16 }, index=oct_nvda_actual.index)
17
18 print(forecast_comparison)

```

	Actual	Forecast
Date		
2024-10-01	117.000000	121.928386
2024-10-02	118.849998	121.536722
2024-10-03	122.849998	121.827043
2024-10-04	124.919998	121.611843
2024-10-07	127.720001	121.771359
2024-10-08	132.889999	121.653118
2024-10-09	132.649994	121.740764
2024-10-10	134.809998	121.675797
2024-10-11	134.800003	121.723954
2024-10-14	138.070007	121.688257
2024-10-15	131.600006	121.714717
2024-10-16	135.720001	121.695104
2024-10-17	136.929993	121.709642
2024-10-18	138.000000	121.698866
2024-10-21	143.710007	121.706854
2024-10-22	143.589996	121.700933
2024-10-23	139.559998	121.705322
2024-10-24	140.410004	121.702068
2024-10-25	141.539993	121.704480
2024-10-28	140.520004	121.702692
2024-10-29	141.250000	121.704017
2024-10-30	139.339996	121.703035

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:
```

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:
```

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:
```

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.


```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:836: ValueWarning:
```

No supported index is available. Prediction results will be given with an integer index beginning at `start`.

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:836: FutureWarning:
```

No supported index is available. In the next version, calling this method in a model without a supported index will result in an exception.

```
In [31]: 1 ### Plot the data
2 plt.figure(figsize=(10, 6))
3 plt.plot(forecast_comparison.index, forecast_comparison['Actual'], l
4 plt.plot(forecast_comparison.index, forecast_comparison['Forecast'],
5
6 # Add labels, title, and legend
7 plt.xlabel('Date')
8 plt.ylabel('Price')
9 plt.title('Actual vs Forecasted NVDA Prices using ARIMA-GARCH in Oct
10 plt.legend()
11
12 # Rotate the x-axis for better readability
13 plt.xticks(rotation=45)
14
15 # Show the plot
16 plt.tight_layout()
17 plt.show()
```



```

In [32]: 1 # Fit ARIMA on AAPL
2 aapl_adj_close = data['Close']['AAPL']
3
4 model_aapl = ARIMA(aapl_adj_close, order=(1, 1, 1))
5 fitted_aapl = model_aapl.fit()
6
7 # Forecast
8 forecast_aapl = fitted_aapl.forecast(steps=len(oct_data))
9
10 oct_aapl_actual = oct_data['Close']['AAPL']
11
12 # Create a DataFrame to compare actual and forecasted prices
13 forecast_comparison = pd.DataFrame({
14     'Actual': oct_aapl_actual,
15     'Forecast': forecast_aapl.values
16 }, index=oct_aapl_actual.index)
17
18 print(forecast_comparison)

```

	Actual	Forecast
Date		
2024-10-01	226.210007	227.803650
2024-10-02	226.779999	227.803582
2024-10-03	225.669998	227.803582
2024-10-04	226.800003	227.803582
2024-10-07	221.690002	227.803582
2024-10-08	225.770004	227.803582
2024-10-09	229.539993	227.803582
2024-10-10	229.039993	227.803582
2024-10-11	227.550003	227.803582
2024-10-14	231.300003	227.803582
2024-10-15	233.850006	227.803582
2024-10-16	231.779999	227.803582
2024-10-17	232.149994	227.803582
2024-10-18	235.000000	227.803582
2024-10-21	236.479996	227.803582
2024-10-22	235.860001	227.803582
2024-10-23	230.759995	227.803582
2024-10-24	230.570007	227.803582
2024-10-25	231.410004	227.803582
2024-10-28	233.399994	227.803582
2024-10-29	233.669998	227.803582
2024-10-30	230.100006	227.803582

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:
```

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:
```

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:
```

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.


```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:836: ValueWarning:
```

No supported index is available. Prediction results will be given with an integer index beginning at `start`.

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:836: FutureWarning:
```

No supported index is available. In the next version, calling this method in a model without a supported index will result in an exception.


```
In [33]: 1 ### Plot the data
2 plt.figure(figsize=(10, 6))
3 plt.plot(forecast_comparison.index, forecast_comparison['Actual'], 1
4 plt.plot(forecast_comparison.index, forecast_comparison['Forecast'],
5
6 # Add labels, title, and legend
7 plt.xlabel('Date')
8 plt.ylabel('Price')
9 plt.title('Actual vs Forecasted AAPL Prices using ARIMA-GARCH in Oct
10 plt.legend()
11
12 # Rotate the x-axis for better readability
13 plt.xticks(rotation=45)
14
15 # Show the plot
16 plt.tight_layout()
17 plt.show()
```



```

In [34]: 1 # Fit ARIMA on MSFT
2 msft_adj_close = data['Close']['MSFT']
3
4 model_msft = ARIMA(msft_adj_close, order=(1, 1, 1))
5 fitted_msft = model_msft.fit()
6
7 # Forecast
8 forecast_msft = fitted_msft.forecast(steps=len(oct_data))
9
10 oct_msft_actual = oct_data['Close']['MSFT']
11
12 # Create a DataFrame to compare actual and forecasted prices
13 forecast_comparison = pd.DataFrame({
14     'Actual': oct_msft_actual,
15     'Forecast': forecast_msft.values
16 }, index=oct_msft_actual.index)
17
18 print(forecast_comparison)

```

	Actual	Forecast
Date		
2024-10-01	420.690002	427.871327
2024-10-02	417.130005	427.930251
2024-10-03	416.540009	427.906896
2024-10-04	416.059998	427.916153
2024-10-07	409.540009	427.912484
2024-10-08	414.709991	427.913938
2024-10-09	417.459991	427.913362
2024-10-10	415.839996	427.913590
2024-10-11	416.320007	427.913500
2024-10-14	419.140015	427.913536
2024-10-15	418.739990	427.913521
2024-10-16	416.119995	427.913527
2024-10-17	416.720001	427.913525
2024-10-18	418.160004	427.913526
2024-10-21	418.779999	427.913525
2024-10-22	427.510010	427.913525
2024-10-23	424.600006	427.913525
2024-10-24	424.730011	427.913525
2024-10-25	428.149994	427.913525
2024-10-28	426.589996	427.913525
2024-10-29	431.950012	427.913525
2024-10-30	432.529999	427.913525

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:
```

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:
```

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:
```

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.


```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:836: ValueWarning:
```

No supported index is available. Prediction results will be given with an integer index beginning at `start`.

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:836: FutureWarning:
```

No supported index is available. In the next version, calling this method in a model without a supported index will result in an exception.

```
In [35]: 1 ### Plot the data
2 plt.figure(figsize=(10, 6))
3 plt.plot(forecast_comparison.index, forecast_comparison['Actual'], l
4 plt.plot(forecast_comparison.index, forecast_comparison['Forecast'],
5
6 # Add labels, title, and legend
7 plt.xlabel('Date')
8 plt.ylabel('Price')
9 plt.title('Actual vs Forecasted MSFT Prices using ARIMA-GARCH in Oct
10 plt.legend()
11
12 # Rotate the x-axis for better readability
13 plt.xticks(rotation=45)
14
15 # Show the plot
16 plt.tight_layout()
17 plt.show()
```



```
In [36]: 1 from pmdarima import auto_arima
2
3 # Auto ARIMA for optimal order selection
4 auto_model = auto_arima(nvda_adj_close, seasonal=False, trace=True)
5 print(auto_model.order)
```

Performing stepwise search to minimize aic

```
ARIMA(2,0,2)(0,0,0)[0] : AIC=384.127, Time=0.02 sec
ARIMA(0,0,0)(0,0,0)[0] : AIC=782.311, Time=0.00 sec
ARIMA(1,0,0)(0,0,0)[0] : AIC=inf, Time=0.01 sec
ARIMA(0,0,1)(0,0,0)[0] : AIC=inf, Time=0.01 sec
ARIMA(1,0,2)(0,0,0)[0] : AIC=381.797, Time=0.01 sec
ARIMA(0,0,2)(0,0,0)[0] : AIC=inf, Time=0.02 sec
ARIMA(1,0,1)(0,0,0)[0] : AIC=383.817, Time=0.01 sec
ARIMA(1,0,3)(0,0,0)[0] : AIC=383.465, Time=0.02 sec
ARIMA(0,0,3)(0,0,0)[0] : AIC=inf, Time=0.04 sec
ARIMA(2,0,1)(0,0,0)[0] : AIC=382.490, Time=0.02 sec
ARIMA(2,0,3)(0,0,0)[0] : AIC=385.611, Time=0.04 sec
ARIMA(1,0,2)(0,0,0)[0] intercept : AIC=373.214, Time=0.06 sec
ARIMA(0,0,2)(0,0,0)[0] intercept : AIC=396.237, Time=0.02 sec
ARIMA(1,0,1)(0,0,0)[0] intercept : AIC=376.350, Time=0.03 sec
ARIMA(2,0,2)(0,0,0)[0] intercept : AIC=372.614, Time=0.07 sec
ARIMA(2,0,1)(0,0,0)[0] intercept : AIC=375.058, Time=0.05 sec
ARIMA(3,0,2)(0,0,0)[0] intercept : AIC=377.397, Time=0.07 sec
ARIMA(2,0,3)(0,0,0)[0] intercept : AIC=374.661, Time=0.09 sec
ARIMA(1,0,3)(0,0,0)[0] intercept : AIC=374.135, Time=0.03 sec
ARIMA(3,0,1)(0,0,0)[0] intercept : AIC=367.317, Time=0.09 sec
ARIMA(3,0,0)(0,0,0)[0] intercept : AIC=372.304, Time=0.02 sec
ARIMA(4,0,1)(0,0,0)[0] intercept : AIC=370.091, Time=0.10 sec
ARIMA(2,0,0)(0,0,0)[0] intercept : AIC=375.668, Time=0.03 sec
ARIMA(4,0,0)(0,0,0)[0] intercept : AIC=373.289, Time=0.02 sec
ARIMA(4,0,2)(0,0,0)[0] intercept : AIC=376.098, Time=0.10 sec
ARIMA(3,0,1)(0,0,0)[0] : AIC=383.555, Time=0.03 sec
```

Best model: ARIMA(3,0,1)(0,0,0)[0] intercept
Total fit time: 1.017 seconds
(3, 0, 1)

```
In [37]: 1 dynamic_forecast = fitted_nvda.get_prediction(start=len(nvda_adj_close),
2                                                    end=len(nvda_adj_close),
3                                                    dynamic=True)
4 forecast_nvda = dynamic_forecast.predicted_mean
```

/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:836: ValueWarning:

No supported index is available. Prediction results will be given with an integer index beginning at `start`.

/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:836: FutureWarning:

No supported index is available. In the next version, calling this method in a model without a supported index will result in an exception.

```
In [38]: 1 from statsmodels.tsa.statespace.sarimax import SARIMAX
2
3 model_sarima = SARIMAX(nvda_adj_close, order=(1, 1, 1), seasonal_order=(1, 1, 1, 1))
4 fitted_sarima = model_sarima.fit()
5 forecast_nvda = fitted_sarima.forecast(steps=len(oct_data))
```

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 4 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 2.60322D+00 |proj g|= 5.62043D-02

At iterate 5 f= 2.52574D+00 |proj g|= 2.80098D-03

/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

This problem is unconstrained.

At iterate 10 f= 2.52560D+00 |proj g|= 1.31536D-04

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
4	11	13	1	0	0	8.294D-06	2.526D+00
F =	2.5255970817443827						

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:836: ValueWarning:
```

No supported index is available. Prediction results will be given with an integer index beginning at `start`.

```
/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:836: FutureWarning:
```

No supported index is available. In the next version, calling this method in a model without a supported index will result in an exception.

STL Decomposition

```
In [39]: 1 from statsmodels.tsa.seasonal import STL
```


```
In [40]: 1 stocks = ["NVDA", "AAPL", "MSFT"]
2 start_date = "2024-07-01"
3 end_date = "2024-09-30"
4 forecast_start_date = "2024-10-01"
5 forecast_end_date = "2024-10-31"
6
7 data = yf.download(stocks, start=start_date, end=end_date)["Adj Clos

[*****100%*****] 3 of 3 completed
```

NVIDIA

```
In [41]: 1 nvda_data = data["NVDA"]
```

```
In [42]: 1 # Ensure the data has a datetime index with a proper frequency
2 nvda_data.index = pd.to_datetime(nvda_data.index)
3
4 # STL decomposition
5 stl = STL(nvda_data, period=5, seasonal=13) # period = 5 to represen
6 result = stl.fit()
7
8 # Plot decomposition
9 result.plot()
10 plt.show()
```



Forecasting

```
In [43]: 1 # Download real test data
2 nvda_test = yf.download(["NVDA"], start=forecast_start_date, end=for

[*****100%*****] 1 of 1 completed
```



```
In [44]: 1 # Forecasting with ARIMA (on the trend component)
2 trend = result.trend.dropna()
3 arima_model = ARIMA(trend, order=(1, 1, 1))
4 arima_fit = arima_model.fit()
5
6 # Forecasting future values
7 forecast_steps = pd.date_range(start=forecast_start_date, end=foreca
8 forecast = arima_fit.get_forecast(len(forecast_steps)).predicted_mea
9 forecast = pd.Series(forecast.values, index=forecast_steps)
```

/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:836: ValueWarning:


No supported index is available. Prediction results will be given with an integer index beginning at `start`.

/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:836: FutureWarning:


No supported index is available. In the next version, calling this method in a model without a supported index will result in an exception.

```
In [46]: 1 # Combine forecasted trend with the seasonal pattern (from the last
2 seasonal_cycle = result.seasonal[-len(forecast):]
3 forecast_with_seasonality = forecast + seasonal_cycle.values
4 forecast_with_seasonality = forecast_with_seasonality.loc[nvda_test.
```

```
In [47]: 1 # Plot predictions
2 plt.figure(figsize=(10, 6))
3 plt.plot(nvda_data, label="Observed (Training)")
4 plt.plot(forecast, label="Forecasted Trend", linestyle="--")
5 plt.plot(forecast_with_seasonality, label="Forecast with Seasonality")
6 plt.legend()
7 plt.title("STL Forecasting for NVDA")
8 plt.show()
```




```
In [48]: 1 # Anomaly detection
2 threshold = 10 # Define acceptable range (e.g., ±10 units)
3 anomalies = abs(nvda_test - forecast_with_seasonality) > threshold
4
5 # Visualize anomalies
6 plt.figure(figsize=(10, 6))
7 plt.plot(nvda_test, label="Real Data")
8 plt.plot(forecast_with_seasonality, label="Forecast", linestyle="--")
9 plt.scatter(nvda_test.index[anomalies], nvda_test[anomalies], color=
10 plt.legend()
11 plt.title("Anomaly Detection in NVDA Data")
12 plt.show()
```



AAPL

```
In [49]: 1 aapl_data = data["AAPL"]
```

```
In [50]: 1 # Ensure the data has a datetime index with a proper frequency
2 aapl_data.index = pd.to_datetime(aapl_data.index)
3
4 # STL decomposition
5 stl = STL(aapl_data, period=5, seasonal=13) # period = 5 to represen
6 result = stl.fit()
7
8 # Plot decomposition
9 result.plot()
10 plt.show()
```



Forecasting

```
In [51]: 1 # Download real test data
2 aapl_test = yf.download(["AAPL"], start=forecast_start_date, end=for

[*****100%*****] 1 of 1 completed
```

```
In [52]: 1 # Forecasting with ARIMA (on the trend component)
2 trend = result.trend.dropna()
3 arima_model = ARIMA(trend, order=(1, 1, 1)) # Adjust ARIMA paramete
4 arima_fit = arima_model.fit()
5
6 # Forecasting future values
7 forecast_steps = pd.date_range(start=forecast_start_date, end=foreca
8 forecast = arima_fit.get_forecast(len(forecast_steps)).predicted_mea
9 forecast = pd.Series(forecast.values, index=forecast_steps)
```

/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/statespace/sarimax.py:978: UserWarning:

Non-invertible starting MA parameters found. Using zeros as starting parameters.

/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:836: ValueWarning:


No supported index is available. Prediction results will be given with an integer index beginning at `start`.

/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:836: FutureWarning:


No supported index is available. In the next version, calling this method in a model without a supported index will result in an exception.

```
In [53]: 1 # Combine forecasted trend with the seasonal pattern (from the last
2 seasonal_cycle = result.seasonal[-len(forecast):]
3 forecast_with_seasonality = forecast + seasonal_cycle.values
4 forecast_with_seasonality = forecast_with_seasonality.loc[aapl_test.
```

```
In [54]: 1 # Plot predictions
2 plt.figure(figsize=(10, 6))
3 plt.plot(aapl_data, label="Observed (Training)")
4 plt.plot(forecast, label="Forecasted Trend", linestyle="--")
5 plt.plot(forecast_with_seasonality, label="Forecast with Seasonality")
6 plt.legend()
7 plt.title("STL Forecasting for AAPL")
8 plt.show()
```




```
In [55]: 1 # Anomaly detection
2 threshold = 10 # Define acceptable range (e.g., ±10 units)
3 anomalies = abs(aapl_test - forecast_with_seasonality) > threshold
4
5 # Visualize anomalies
6 plt.figure(figsize=(10, 6))
7 plt.plot(aapl_test, label="Real Data")
8 plt.plot(forecast_with_seasonality, label="Forecast", linestyle="--")
9 plt.scatter(aapl_test.index[anomalies], aapl_test[anomalies], color=
10 plt.legend()
11 plt.title("Anomaly Detection in AAPL Data")
12 plt.show()
```



Microsoft

```
In [56]: 1 msft_data = data["MSFT"]
```

```
In [57]: 1 # Ensure the data has a datetime index with a proper frequency
2 msft_data.index = pd.to_datetime(msft_data.index)
3
4 # STL decomposition
5 stl = STL(msft_data, period=5, seasonal=13) # period = 5 to represen
6 result = stl.fit()
7
8 # Plot decomposition
9 result.plot()
10 plt.show()
```



Forecasting

```
In [58]: 1 # Download real test data
2 msft_test = yf.download(["MSFT"], start=forecast_start_date, end=for

[*****100%*****] 1 of 1 completed
```



```
In [59]: 1 # Forecasting with ARIMA (on the trend component)
2 trend = result.trend.dropna()
3 arima_model = ARIMA(trend, order=(1, 1, 1)) # Adjust ARIMA paramete
4 arima_fit = arima_model.fit()
5
6 # Forecasting future values
7 forecast_steps = pd.date_range(start=forecast_start_date, end=foreca
8 forecast = arima_fit.get_forecast(len(forecast_steps)).predicted_mea
9 forecast = pd.Series(forecast.values, index=forecast_steps)
```

/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning:

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:836: ValueWarning:


No supported index is available. Prediction results will be given with an integer index beginning at `start`.

/Users/cynthiadu/anaconda3/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:836: FutureWarning:

No supported index is available. In the next version, calling this method in a model without a supported index will result in an exception.

```
In [60]: 1 # Combine forecasted trend with the seasonal pattern (from the last
2 seasonal_cycle = result.seasonal[-len(forecast):]
3 forecast_with_seasonality = forecast + seasonal_cycle.values
4 forecast_with_seasonality = forecast_with_seasonality.loc[msft_test.]
```

```
In [61]: 1 # Plot predictions
2 plt.figure(figsize=(10, 6))
3 plt.plot(msft_data, label="Observed (Training)")
4 plt.plot(forecast, label="Forecasted Trend", linestyle="--")
5 plt.plot(forecast_with_seasonality, label="Forecast with Seasonality")
6 plt.legend()
7 plt.title("STL Forecasting for MSFT")
8 plt.show()
```



```
In [62]: 1 # Anomaly detection
2 threshold = 10 # Define acceptable range (e.g., ±10 units)
3 anomalies = abs(msft_test - forecast_with_seasonality) > threshold
4
5 # Visualize anomalies
6 plt.figure(figsize=(10, 6))
7 plt.plot(msft_test, label="Real Data")
8 plt.plot(forecast_with_seasonality, label="Forecast", linestyle="--")
9 plt.scatter(msft_test.index[anomalies], msft_test[anomalies], color=
10 plt.legend()
11 plt.title("Anomaly Detection in MSFT Data")
12 plt.show()
```

